

Cloud and Cluster Data Management

SCALABLE CONSISTENCY AND TRANSACTION MODELS

THANKS TO M. GROSSNIKLAUS

Outline

- Sharding & Replication –
 - What are they?
 - What issues do they introduce?
- CAP Theorem (Consistency, Availability, Partition-Tolerance – can't have it all)
 - Also look at consistency vs. latency
- Eventual Consistency
 - What is it?
 - What are different models of eventual consistency?
 - Also look at configurations of readers and writers for replication and consistency
- Vector clocks: causal consistency

Sharding and Replication

- Sharding (Partitioning)
 - Breaking a database into several collections (*shards*)
 - Each data item (e.g., a document) goes in one shard
- Replication
 - Have multiple copies of a database
 - Each data item lives in several places

Can combine sharding and replication

Why do systems shard and replicate?

Issues with Sharding and Replication

Sharding: If an operation needs multiple data items, it might need to access several shards

Replication: Trying to keep replicas in sync with each other.


Other issues:

- Replications requires more storage space
- Time of access for both sharding and replications
- If more machines are being added, how would data be redistributed?
- How to locate data among nodes for sharding

Managing Replicas 1

Strategy 1: Write all


- Write: Update both, synchronously
- Read: Access either



Managing Replicas 2

Strategy 2: Write one

- Write: Update one, propagate update asynchronously
- Read: Access either (or both?)



Brewer's Conjecture

- Three properties that are desirable and expected from real-world shared-data systems
 - **C**: data consistency – you get most recent write or an error
 - **A**: availability – every request gets a (non-error) response
 - **P**: tolerance of network partition
- At *PODC 2000* (Portland, OR), Eric Brewer made the conjecture that only two of these properties can be satisfied by a system at any given time
- Conjecture was formalized and confirmed by MIT researchers Seth Gilbert and Nancy Lynch in 2002
- Now known as the **CAP Theorem**

Data Consistency

- Database systems typically implement ACID transactions
 - **Atomicity**: “all or nothing”
 - **Consistency**: transactions never observe or result in inconsistent data
 - **Isolation**: transactions are not aware of concurrent transactions
 - **Durability**: once committed, the state of a transaction is permanent
- Useful in automated business applications
 - banking: at the end of a transaction the sum of money in both accounts is the same as before the transaction
 - online auctions: the last bidder wins the auction
- There are applications that can deal with looser consistency guarantees and periods of inconsistency

Availability

- Services are expected to be highly available
 - every request should receive a response
 - if you can read a data item, you can update it
 - it can create real-world problems when a service goes down
- Realistic goal
 - service should be as available as the network it runs on
 - if any instance of a service on the network is available, the service should be available

Partition-Tolerance


- A service should continue to perform as expected
 - if some nodes crash
 - if some communication links fail
- One desirable fault tolerance property is resilience to a network partitioning into multiple components
- In cloud computing, node and communication failures are not the exception but everyday events

Managing Replicas: Activity

Strategy 3: Primary-Secondary

1. Write: Update primary, propagate to secondary asynchronously
2. Read Strong: Access primary
3. Read Weak: Access either

For each op, does it give consistency or availability in a partition?



Participation Question 1

In the primary-secondary strategy, for each of the following operations, does it give consistency or availability in a partition?

1. Write: Update primary, propagate to secondary asynchronously
2. Read Strong: Access primary
3. Read Weak: Access either

