

Cloud and Cluster Data Management

SCALABLE CONSISTENCY AND TRANSACTION MODELS

PART 3

Outline

- Sharding & Replication –
 - What are they?
 - What issues do they introduce?
- CAP Theorem (Consistency, Availability, Partition-Tolerance – can't have it all)
 - Also look at consistency vs. latency
- Eventual Consistency
 - What is it?
 - What are different models of eventual consistency?
 - Also look at configurations of readers and writers for replication and consistency
- Vector clocks: causal consistency

Configurations

- Definitions
 - **N**: number of nodes that store a replica
 - **W**: number of replicas that need to acknowledge a write operation
 - **R**: number of replicas that are accessed for a read operation
- $W+R > N$
 - e.g. **synchronous replication** ($N=3$, $W=2$, and $R=2$)
 - write set and read set always overlap
 - strong consistency can be guaranteed through **quorum protocols**
 - risk of reduced availability: in basic quorum protocols, operations fail if fewer than the required number of nodes respond, due to node failure
- $W+R = N$
 - e.g. **asynchronous replication** ($N=2$, $W=1$, and $R=1$)
 - strong consistency cannot be guaranteed

Configurations

- $R=1, W=N$
 - optimized for **read access**: single read will return a value
 - write operation involves all nodes and risks not succeeding
- $R=N, W=1$
 - optimized for **write access**: write operation involves only one node and relies on asynchronous updates to other replicas
 - read operation involves all nodes and returns “latest” value
 - durability is not guaranteed in presence of failures
- $W < (N+1)/2$
 - risk of conflicting writes
- $W+R \leq N$
 - **weak/eventual consistency**

Participation Q1

Your team is designing a **global e-commerce database** with:

- **High read volume** (customers browsing products).
- **Critical financial transactions** (orders/payments must be consistent).
- **Multiple data centers** across different regions.

You must choose a **replication model**:

1. **Strong Consistency ($W+R > N$)** – Guarantees latest data, but increases latency.
2. **Eventual Consistency ($W+R \leq N$)** – Faster but may return stale data.

Team Roles:

- **Database Engineer:** Focus on consistency & replication.
- **App Developer:** Focus on read/write performance.
- **Business Manager:** Focus on availability & cost.

Discussion Questions (5 min):

1. Which model does your team choose and why?
2. What is the biggest risk of your choice?
3. How would you adjust W and R to balance performance & consistency?

BASE

- **B**asically **A**vailable, **S**oft state, **E**ventually Consistent
- As consistency is achieved eventually, conflicts have to be resolved at some point
 - read repair
 - write repair
 - asynchronous repair
- Conflict resolution is typically based on a global (partial) ordering of operations that (deterministically) guarantees that all replicas resolve conflicts in the same way
 - client-specified timestamps
 - vector clocks

Vector Clocks

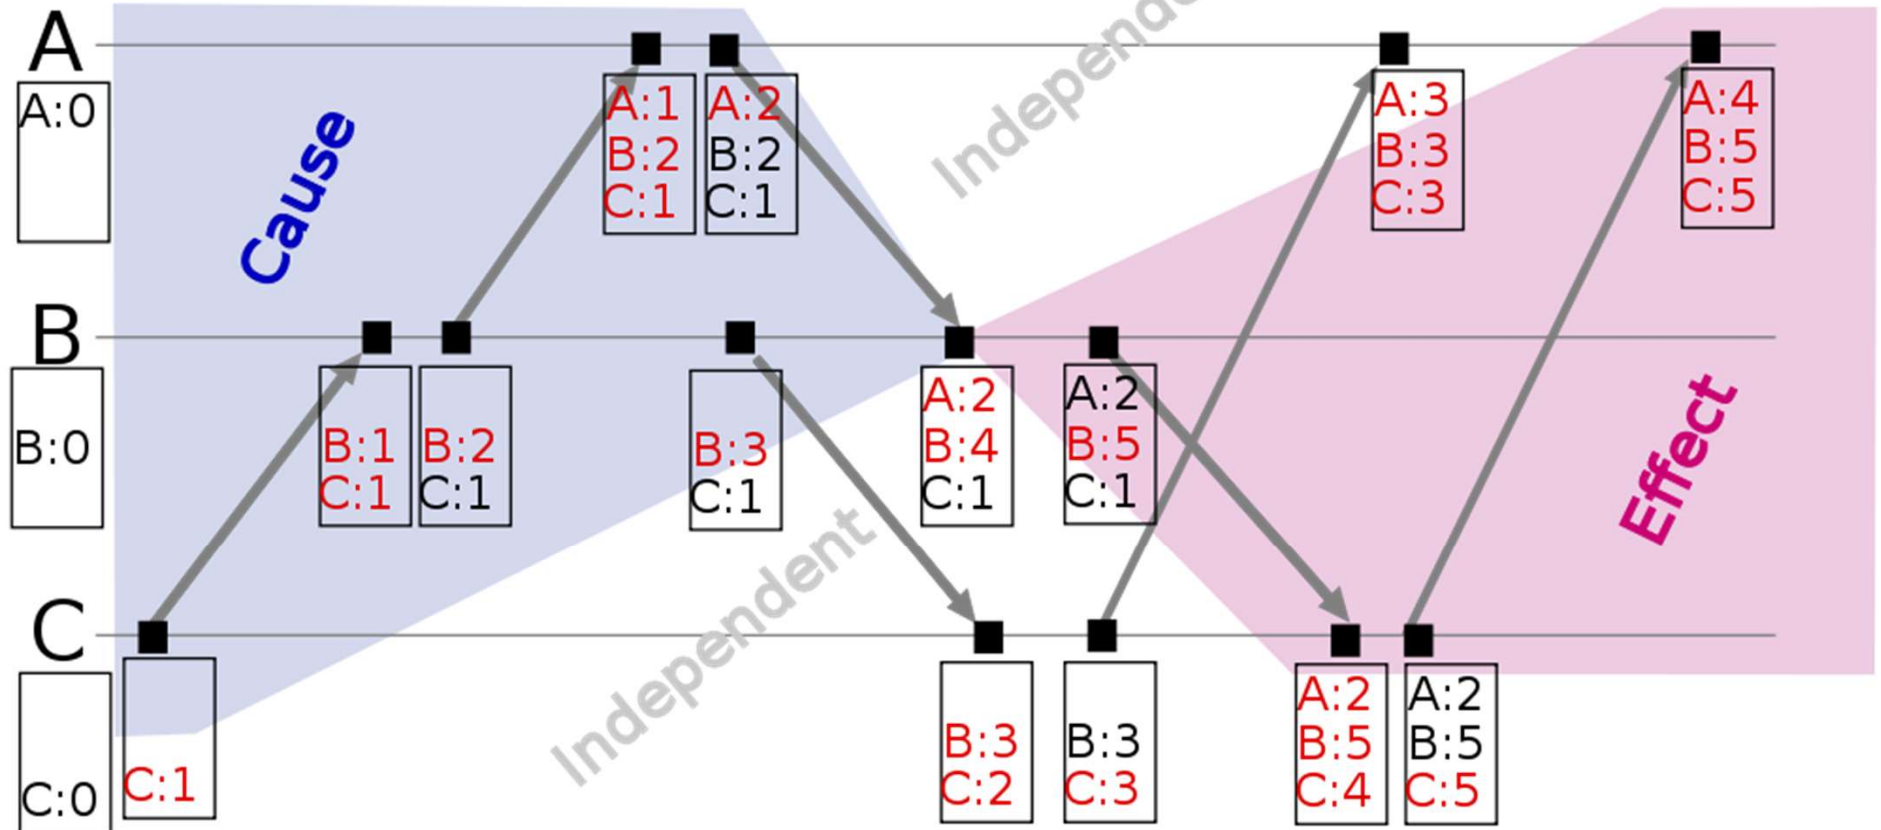
- Generate a partial ordering of events in a distributed system and detecting causality violations
- A **vector clock** of a system of n processes is an vector of n logical clocks (one clock per process)
 - messages contain the state of the sending process's logical clock
 - a local “smallest possible values” copy of the global vector clock is kept in each process
- Vector clocks algorithm was independently developed by Colin Fidge and Friedemann Mattern in 1988

Update Rules for Vector Clocks

- All clocks are initialized to zero
- A process increments its own logical clock in the vector by one
 - each time it experiences an internal event
 - each time a process prepares to send a message
 - each time a process receives a message
- Each time a process sends a message, it transmits the entire vector clock along with the message being sent
- Each time a process receives a message, it updates each element in its vector by taking the pair-wise maximum of the value in its own vector clock and the value in the vector in the received message

Vector Clock Example

Time →



Participation Question 2

Detecting Causality Violations

- Your team is managing a distributed chat application where messages are replicated across multiple servers using vector clocks. Consider the following event sequence:
 1. Alice sends a message: **"Let's meet at 3 PM"** (Vector clock: **[1,0,0]**).
 2. Bob replies: **"Okay, see you then"** (Vector clock: **[1,1,0]**).
 3. Alice, unaware of Bob's response, updates the time: **"Actually, let's meet at 4 PM"** (Vector clock: **[2,0,0]**).
 4. Carol receives both messages at the same time.

Discussion Questions

1. Using vector clocks, how can Carol determine if these messages are causally related or concurrent?
2. Which message, if any, happened before the other, and how do you know?
3. If Alice had seen Bob's response before updating the meeting time, how would that have changed her vector clock?

References

- S. Gilbert and N. Lynch: **Brewer's Conjecture and the Feasibility of Consistent, Available and Partition-Tolerant Web Services.** *SIGACT News* 33(2), pp. 51-59, 2002.
- D. J. Abadi: **Consistency Tradeoffs in Modern Distributed Database System Design: CAP is Only Part of the Story.** *Computer*, vol. 45, no. 2, pp. 37-42, Feb. 2012.
- A. Thomson and D. J. Abadi. **The Case for Determinism in Database Systems.** *Proc. VLDB Endow.* 3, 1-2 (September 2010).
- D. Terry. **Replicated Data Consistency Explained Through Baseball.** *Commun. ACM* 56, 12, Dec. 2013.
- W. Vogels: **Eventually Consistent.** *ACM Queue* 6(6), pp. 14-19, 2008.