

Data Models

Key/Value: Amazon Dynamo


Part 2

Thanks to Dave Maier, M. Grossniklaus & K. Tufte

Replication

- Replication is used to achieve high availability and durability
- Every data item is replicated at N hosts
 - N is configured “per-instance” of Dynamo
 - each key k is assigned a **coordinator** host that handles write requests for k
 - coordinator is also in charge of replication of data items within its range
- Algorithm
 1. coordinator stores data item with key k locally
 2. coordinator stores data item at $N-1$ clockwise successors nodes
- Every node is responsible for the region of the ring between itself and its N^{th} predecessor
- Preference list
 - enumerates nodes that are responsible for storing a key k
 - contains more than N nodes to account for node failures

Replication

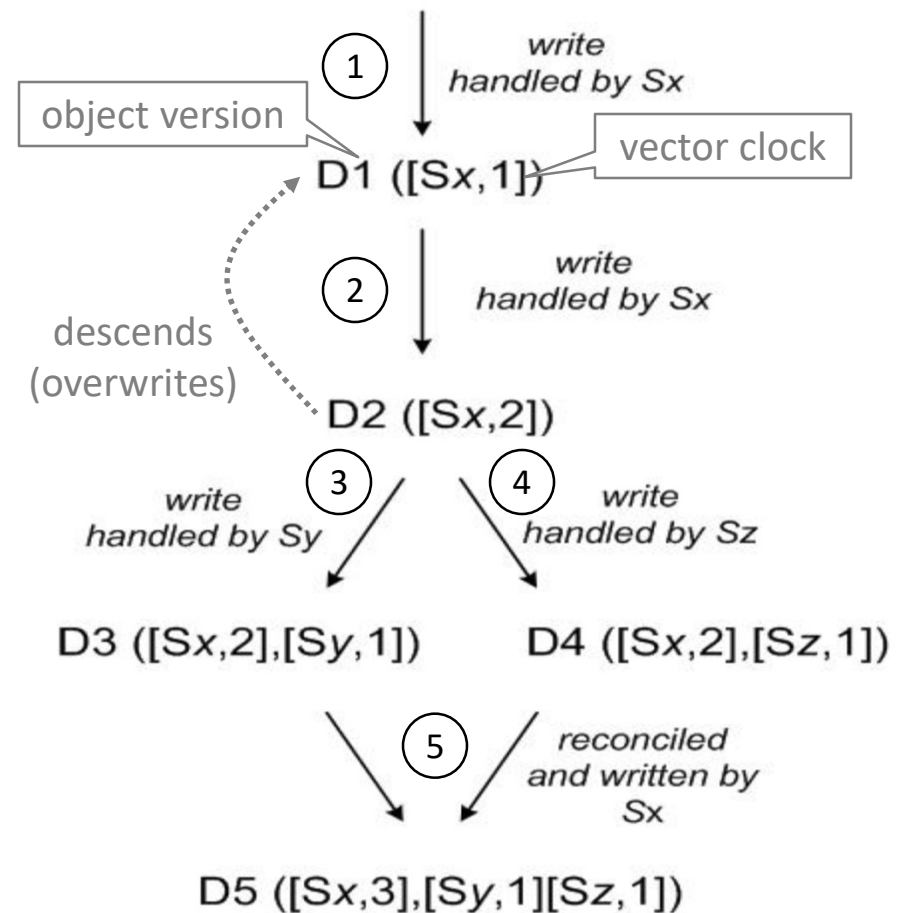


Data Versioning

- Eventual consistency by asynchronous updates of replicas
 - put() may return to caller before all replicas have been updated
 - subsequent get() may return objects that have not been updated yet
- “Always writable” design
 - result of each modification is a new and immutable version of the data
 - multiple versions may exist in the system at the same time
 - vector clocks capture causality between different versions of an object
- Version reconciliation
 - **system-based**: most versions simply subsume the previous version
 - **client-based**: if failures combined with concurrent updates lead to branches, the client (app) needs to collapse multiple branches into one
- Context passed between get and put operations contains vector clock information

Data Versioning Evolution

1. Client A
 - writes new object D
 - node S_x writes version D1
2. Client A
 - updates object D
 - node S_x writes version D2
3. Client A
 - updates object D
 - node S_y writes version D3
4. Client B
 - reads and updates object D
 - node S_z writes version D4
5. Client C
 - reads object D (i.e., D3 and D4)
 - client performs reconciliation
 - node S_x writes version D5



PQ1: How Would You Reconcile Two Carts?

<u>Item</u>	<u>Qty</u>	<u>Price-Each</u>	<u>Item</u>	<u>Qty</u>	<u>Price-Each</u>
Mixer	1	199.00	Mixer	1	199.00
Tea	3	7.50	Tea	2	7.50
Teapot	1	23.25	Spatula	4	3.30
Spatula	4	3.30	Strainer	2	12.80
Strainer	1	12.80	Kettle	1	22.65


What was your general strategy?

Execution of Get and Put Operations

- **Coordinator**: node that handles read or write operation
 - typically the first of the top N nodes of the preference list for a write
 - requests that hit a node that is not in the top N of the requested key's preference list are forwarded to the appropriate coordinator
- Quorum-like Protocol
 - R : minimum number of nodes that must participate in successful read
 - W : minimum number of nodes that must participate in successful write
- “Sloppy Quorum”
 - read and write operations are performed on the first N **healthy** nodes
 - not guaranteed be the first N nodes encountered while walking the ring
- Hinted Handoff
 - if a node is unreachable, a **hinted** replica is sent to next healthy node
 - nodes receive hinted replicas keep them in a separate database
 - hinted replica is delivered to original node when it recovers


Execution of Get Operation

$N = 3$
 $R = 2$
 $W = 2$



Execution of Put Operation

$N = 3$
 $R = 2$
 $W = 2$

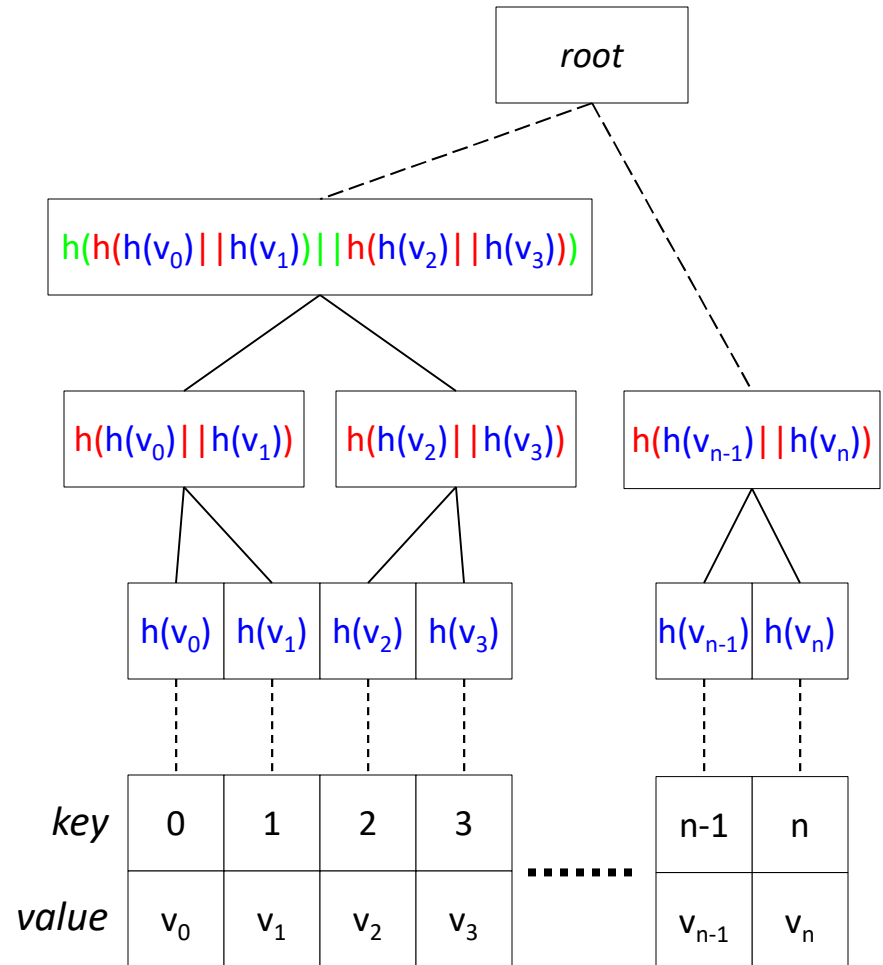


Replica Synchronization

- Scenarios exist where hinted replicas become unavailable before being returned to the original node
 - Dynamo implements an anti-entropy (replica synchronization) protocol
 - Dynamo uses Merkle trees to detect inconsistencies between replicas
- Algorithm
 - every physical node maintains a separate Merkle tree for each hosted key range, i.e., the set of keys covered by a virtual host
 - to check if their key ranges are up-to-date, two nodes exchange the roots of the Merkle tree of the key ranges they have in common
 - if the value of the roots are equal, the key ranges are up-to-date
 - if not, they recursively exchange the values of the children until they reach the leaves of the Merkle tree
 - at that point, the inconsistent keys are identified and the corresponding values can be exchanged

Merkle Tree


- Hash tree
 - **leaves** are hashes of the values of individual keys
 - **non-leaves** are hashes of their respective children nodes
- Properties
 - efficient verification of large data structures
 - every branch can be processed independently
- Applications
 - ZFS file systems
 - Git revision control system
 - NoSQL systems (e.g., Dynamo, Cassandra, and Riak)



Membership and Failure Detection

- Explicit mechanism to initiate addition and removal of nodes
 - outage or failure rarely signifies permanent departure of node
 - manual error could result in unintentional startup of node
 - these events should not result in rebalancing of partition assignment or repair of unreachable replica
- Full membership model
 - administrative command issued to join/remove a node to/from ring
 - node that handles request updates its persistent membership table
 - gossip-based protocol propagates membership changes
- Failure detection
 - communication failures avoided based on a purely local failure notion
 - globally consistent view on failure state is not required thanks to explicit join/remove commands

Full Membership Model



PQ2: Membership Management and Failure Detection

Your team is managing a **distributed online ticket booking system** that relies on DynamoDB.

A **server failure** has been detected, and you must decide how to handle **membership changes and failure recovery**.

Discuss the following:

- 1. Why does Dynamo use an explicit join/remove command instead of automatic failure detection for managing membership?**
- 2. What advantages does a gossip-based protocol offer for updating membership across nodes?**
- 3. How does hinted handoff help maintain availability during temporary node failures? What happens if the hinted replica is lost?**

Prepare to present your team's insights.

Implementation

Three main software components

1. Local persistence engine

- plug-in architecture supports different storage engines
- storage engine chosen based on application's object size distribution
- *BerkeleyDB*: objects in the order of tens of kilobytes
- *MySQL*: larger objects

2. Request coordination

- built on top of an event-driven messaging infrastructure
- communication implemented using Java NIO channels
- client requests create a state machine on node that received request
- each state machine handles exactly one client request

3. Membership and failure detection

Configurability

N	R	W	Application
3	2	2	Consistent, durable, interactive user state (typical configuration)
n	1	n	High-performance read engine
1	1	1	Distributed web cache

Typical Usage Patterns

- Business logic-specific reconciliation
 - each data object is replicated across multiple nodes
 - client applications performs reconciliation in case of divergent versions
 - **example**: merging different versions of a customer's shopping cart
- Timestamp-based reconciliation
 - Dynamo performs simple timestamp-based reconciliation
 - “last write wins”, i.e., object with largest physical times is selected
 - **example**: maintaining a customer's session information
- High-performance read engine
 - $R = 1, W = N$
 - high read request rate with only a small number of updates
 - **example**: management of product catalog and promotional items